



The AI/Data Science Professional

Responsible AI/DS



Acknowledgement of Country

The AI/Data Science Professional
Course Coordinator: Flora Salim

—
What's next...

RMIT University acknowledges the people of the Woi wurrung and Boon wurrung language groups of the eastern Kulin Nation on whose unceded lands we conduct the business of the University.

RMIT University respectfully acknowledges their Ancestors and Elders, past and present. RMIT also acknowledges the Traditional Custodians and their Ancestors of the lands and waters across Australia where we conduct our business.



Ngarara Place



The AI/Data Science Professional – Week 2: Responsible AI/DS

Responsible AI/Data Science



Objectives

- Passion about AI / Data Science :)
- Ethical Mindset
 - ✧ Critical thinking
 - ✧ Research
 - ✧ Evaluation
- Code of Practice
- Technologies
- Professional practice
 - ✧ Communication
 - ✧ Team work

An AI/DS professional



Course Learning Outcomes



- CLO1 - Standards, ethical and social considerations
- CLO2 - Analyse and discuss impact
- CLO3 - Case studies
- CLO4 - Communication
- CLO5 - Research Principles

Refer to Course Guide for the actual description of the learning outcomes.

Responsible AI/DS

- Cost and Risk
 - ✧ Privacy
 - ✧ Fairness
 - ✧ Transparency and Explainability
 - ✧ Accountability, Responsibility and Governance
 - ✧ Human Right and Social Wellbeing
- Existing AI ethics and governance framework
- What to do this week
- Guest talk on Research Integrity and Ethics

AI adoption in industry

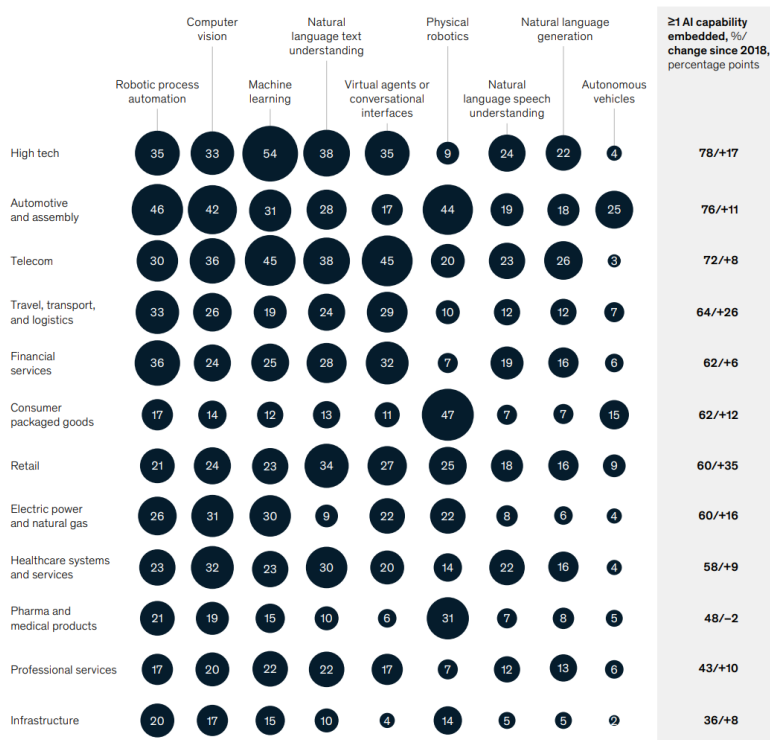


AI global survey by McKinsey

<https://www.mckinsey.com/featured-insights/artificial-intelligence/global-ai-survey-ai-proves-its-worth-but-few-scale-impact>

High tech leads in AI adoption, and industries are generally using the AI capabilities most relevant to their value chains.

Organizations' AI capabilities,¹ % of respondents² by industry



¹Embedded in ≥1 product and/or business process for ≥1 function or business unit.

²Respondents who said "don't know" or "none of the above" are not shown. For high tech, n = 277; for automotive and assembly, n = 128; for telecom, n = 93; for travel, transport, and logistics, n = 83; for financial services, n = 396; for consumer packaged goods, n = 72; for retail, n = 94; for electric power and natural gas, n = 82; for healthcare systems and services, n = 78; for pharma and medical products, n = 96; for professional services, n = 331; and for infrastructure, n = 91.

AI adoption leads to increased revenue and decreased cost




AI global survey by McKinsey

<https://www.mckinsey.com/featured-insights/artificial-intelligence/global-ai-survey-ai-proves-its-worth-but-few-scale-impact>

Revenue increases from adopting AI are reported most often in marketing and sales, and cost decreases most often in manufacturing.

Cost decrease and revenue increase from AI adoption, by function,¹ % of respondents²



¹Marketing and sales includes the following use cases: customer-service analytics, customer segmentation, channel management, prediction of likelihood to buy, pricing and promotion, closed-loop marketing, marketing-budget allocation, churn reduction, and next product to buy. For product and service development: product-feature optimization, product-development-cycle optimization, creation of new AI-based enhancements, and creation of new AI-based products. For supply-chain management: logistics-network optimization, sales and parts forecasting, warehouse optimization, inventory and parts optimization, spend analytics, and sales and demand forecasting. For manufacturing: predictive maintenance and yield, energy, and throughput optimization. For service operations: service-operations optimization, contact-center automation, and predictive service and intervention. For strategy and corporate finance: capital allocation, treasury management, and M&A support. For risk: risk modeling/analytics, and fraud/debt analytics. For HR: performance management and organization-design, workforce-deployment, and talent-management optimization.

²Question asked only of respondents who said their companies adopted AI in given use case. Figures were calculated after removing respondents who said "don't know" or "not applicable; we are not tracking revenue related to AI"; respondents who said "no change" are not shown.

Human Rights and Social Wellbeing



Image sourced from The Intercept Source
<https://www.flickr.com/photos/itupictures/39621400230>



[Discussion]

Are you concerned about the
risk of AI making us jobless?



Image sourced from The Intercept Source from <https://chatbotlife.com/will-ai-cause-mass-unemployment-f26537e5f25a>

AI risks

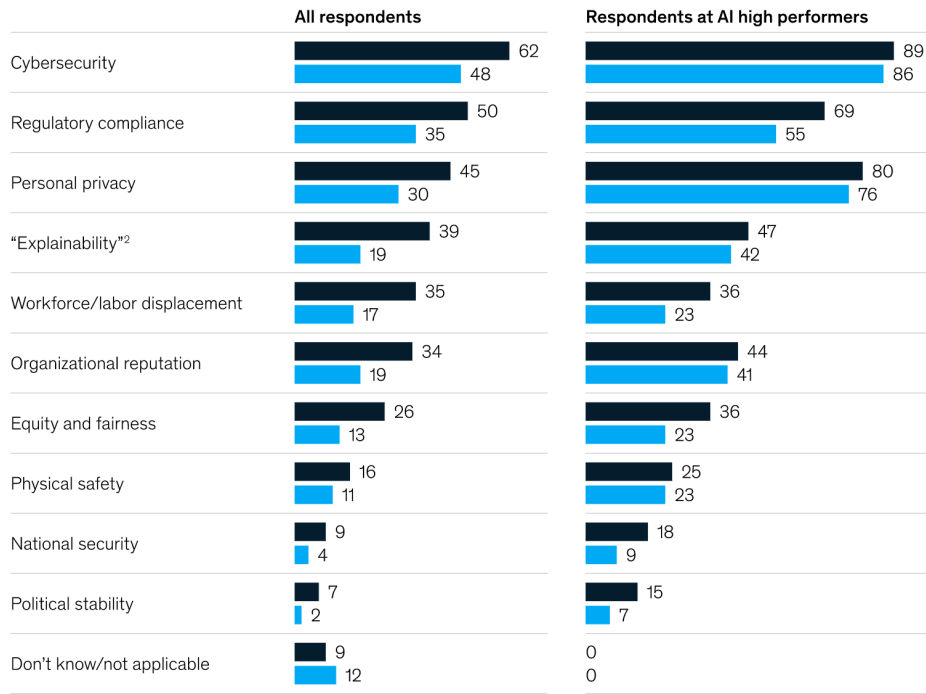


AI global survey by McKinsey

<https://www.mckinsey.com/featured-insights/artificial-intelligence/global-ai-survey-ai-proves-its-worth-but-few-scale-impact>

Risks that organizations consider relevant and are working to mitigate,
% of respondents¹

■ Relevant risk
■ Mitigated risk



¹Question asked only of respondents who said their companies had embedded or piloted ≥1 AI capability; n = 1,872.

²Ability to explain how AI models come to their decisions.

Job loss/changes due to AI


AI global survey by McKinsey

<https://www.mckinsey.com/featured-insights/artificial-intelligence/global-ai-survey-ai-proves-its-worth-but-few-scale-impact>

Respondents in automotive and telecom report the deepest AI-related workforce cuts to date and predict the most going forward.

Change in workforce due to AI adoption, % of respondents

■ Decrease >10% ■ Decrease 3–10% ■ Increase 3–10% ■ Increase >10%



¹Change in workforce in past year because of AI adoption. Question only asked of respondents who say their companies have piloted or embedded 21 AI capability. Respondents who said "don't know" are not shown. For automotive and assembly, n = 111; for telecom, n = 81; for infrastructure, n = 63; for retail, n = 74; for financial services, n = 333; for professional services, n = 235; for high tech, n = 246; for pharma and medical products, n = 71; for consumer packaged goods, n = 55; for travel, transport, and logistics, n = 69; for healthcare systems and services, n = 60.

²A decrease or increase of 52%.

³Expected change in workforce in next 3 years because of AI adoption, relative to size if AI had not been adopted. Question was asked only of respondents who say their companies have piloted or embedded 21 AI capability, or plan to do so in the next 3 years. Respondents who said "don't know" are not shown. For automotive and assembly, n = 113; for telecom, n = 85; for financial services, n = 341; for professional services, n = 245; for high tech, n = 253; for pharma and medical products, n = 78; for consumer packaged goods, n = 58; for travel, transport, and logistics, n = 70; and for healthcare systems and services, n = 67.

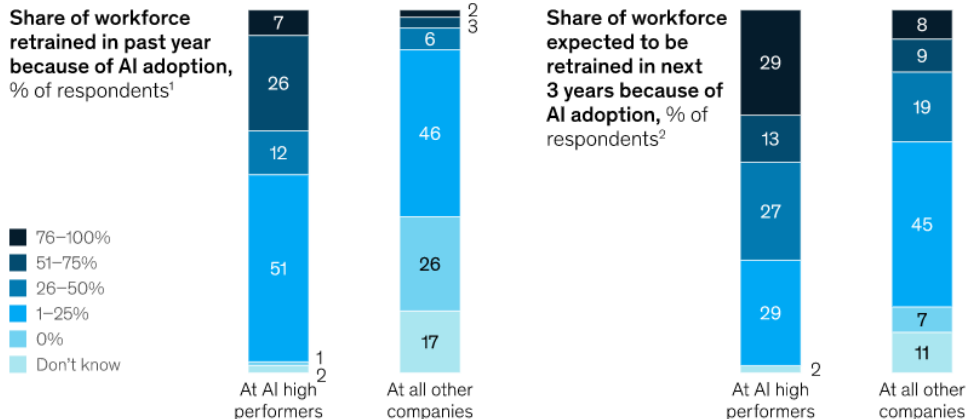
Workforce retraining due to AI



AI global survey by McKinsey

<https://www.mckinsey.com/featured-insights/artificial-intelligence/global-ai-survey-ai-proves-its-worth-but-few-scale-impact>

Respondents at high performers report larger retraining efforts as a result of AI than others do.



Note: Figures may not sum to 100%, because of rounding.

¹Question was asked only of respondents who say their companies have piloted or embedded one or more AI capabilities. For respondents at high performers, n = 54; for all others, n = 1,818.

²Question was asked only of respondents who say their companies have piloted or embedded one or more AI capabilities, or plan to do so in the next three years. For respondents at high performers, n = 54; for all others, n = 1,892.

McKinsey
& Company

A case study: virtual assistant



What are the benefits and the risks?

Google Duplex

<https://youtu.be/D5VN56jQMWM>

Jarvis (by Mark Zuckerberg)

<https://youtu.be/vvimBPJ3XGQ>

AI Ethics Principles



Human, social
and
environmental
wellbeing

Human-centred
values

Fairness

Privacy
protection and
security

Reliability and
safety

Transparency
and
explainability

Contestability

Accountability

Source from <https://www.industry.gov.au/data-and-publications/building-australias-artificial-intelligence-capability/ai-ethics-framework/ai-ethics-principles>

Fairness

“AI has the potential to help humans make fairer decisions—but only if we carefully work toward fairness in AI systems as well.” -- by Jake Silberg and James Manyika





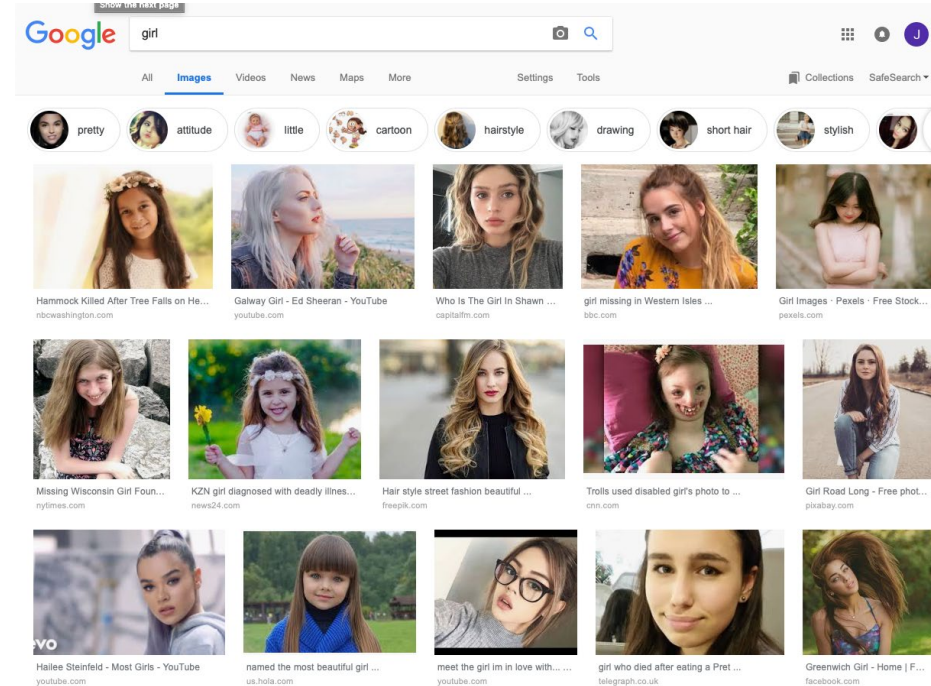
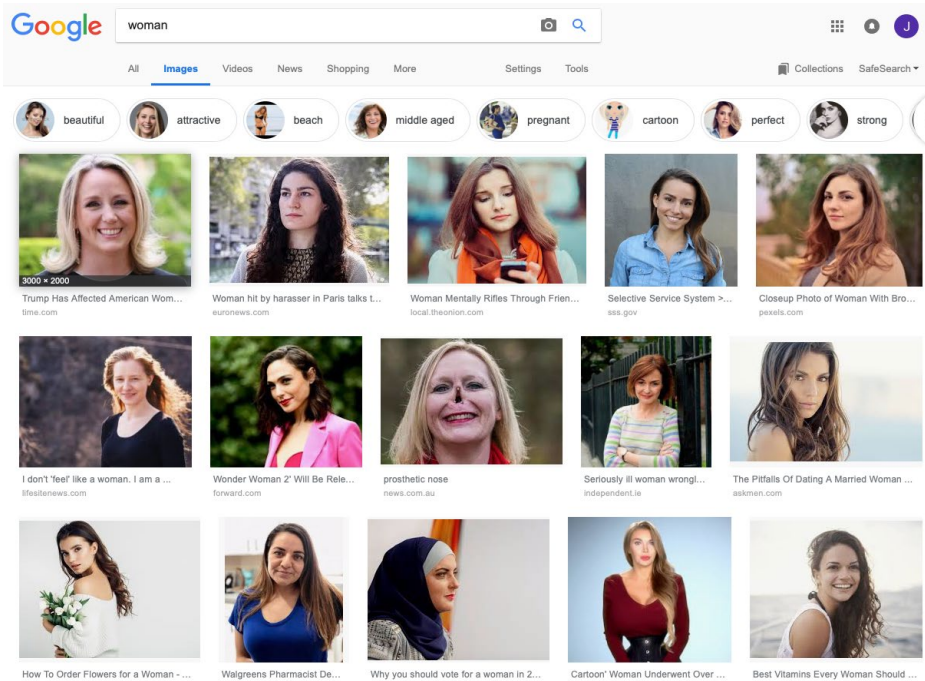
Image sourced from <https://medium.com/thoughts-and-reflections/racial-bias-and-gender-bias-examples-in-ai-systems-7211e4c166a1>



[Discussion]

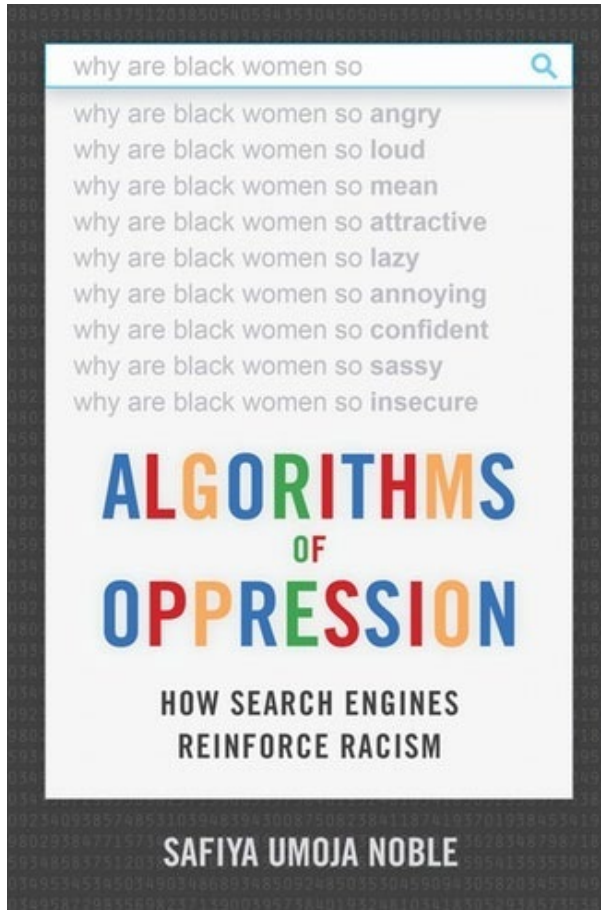
Is there any AI/DS application that concern you with regards to Fairness?

Search Engine Bias



Source from <https://scroll.in/article/921305/google-search-results-reflect-its-algorithms-bias-against-women-and-people-of-colour>


Suggestive autocompletion



How biased are our algorithms? |
Safiya Umoja Noble | TEDxUIUC

Other examples

- PredPol
- Facial recognition
- Google images search for “CEO”
- Facebook falsely feeds the intifada



Transparency and Explainability



Image sourced from
<https://www.nvidia.com/>



[Discussion]


When do you think self-driving cars will be widespread in Australia? And why?



[Discussion]

Are there any other AI applications that concern you in regards to Transparency and Explainability?

Privacy Concern



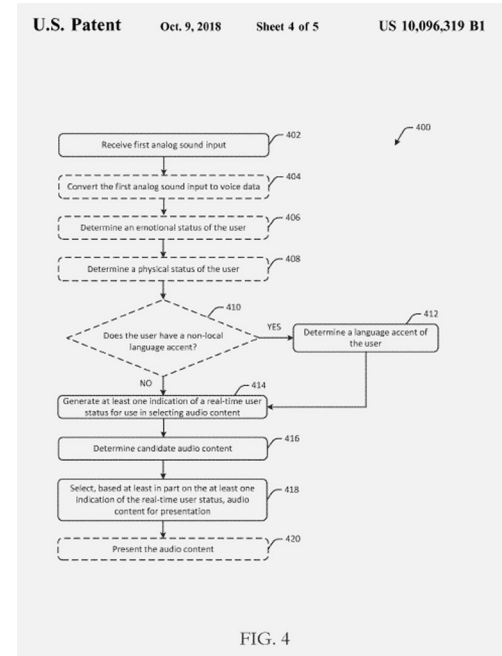
Source from https://theintercept.imgix.net/wp-uploads/sites/1/2018/11/Intercept_Echo_v2-3.5MB-2-1542062294.gif

Case study: AMAZON'S ACCENT RECOGNITION TECHNOLOGY



AT THE BEGINNING of October 2018, Amazon was quietly issued a patent that would allow its virtual assistant Alexa to decipher a user's physical characteristics and emotional state based on their voice. Characteristics, or “voice features,” like language accent, ethnic origin, emotion, gender, age, and background noise would be immediately extracted and tagged to the user's data file to help deliver more targeted advertising.

The algorithm would also consider a customer's physical location — based on their IP address, primary shipping address, and browser settings — to help determine their accent.



The illustration from Document:
United States Patent and
Trademark Office



[Discussion]

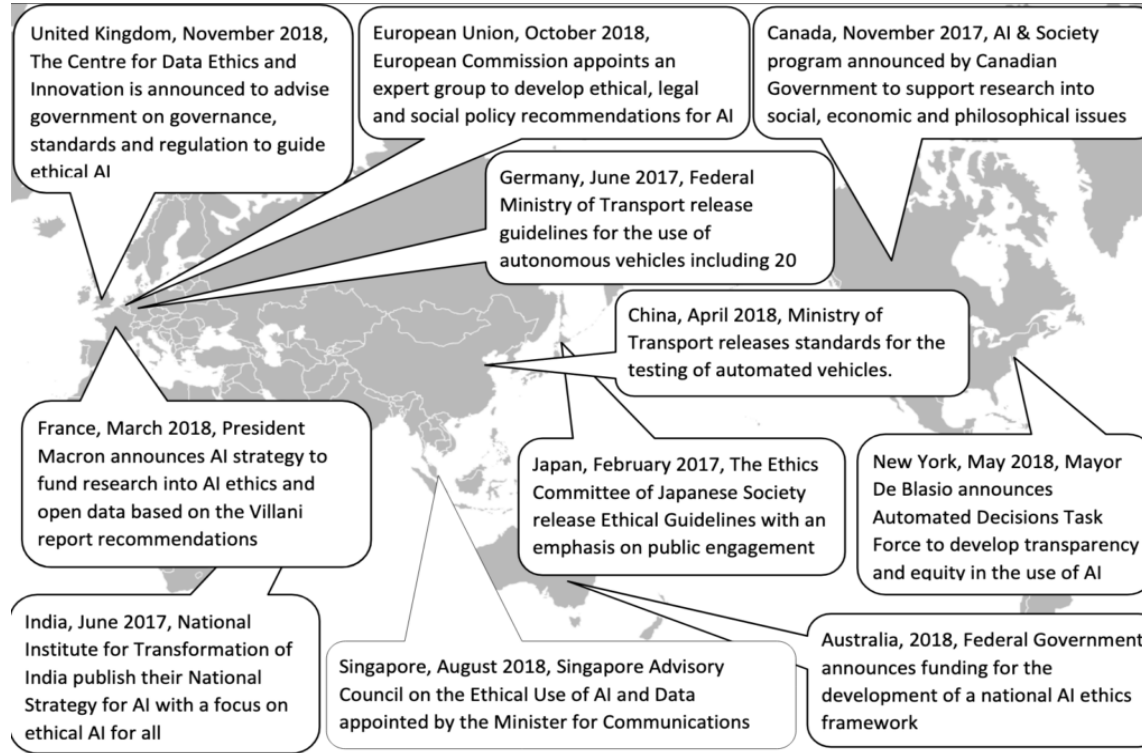
Is there any other AI/DS
application that concern you
with regards to Privacy?

Accountability, Responsibility and Governance



<https://medium.com/@viriniadignum/the-art-of-ai-accountability-responsibility-transparency-48666ec92ea5>

Existing AI ethics and governance framework



Australian Framework



- Government and automated decisions
- Australia's international human rights obligations and anti-discrimination legislation
- Data-sharing legislation in Australia
- Privacy Act

https://consult.industry.gov.au/strategic-policy/artificial-intelligence-ethics-framework/supporting_documents/ArtificialIntelligenceethicsframeworkdiscussionpaper.pdf

Further discussion



- Week 3: working with data & users + *Guest talk* by Amazon
- Week 4: FATE in AI (continued)
- Week 5: Privacy & security
- Week 6: AI Bias and Fairness
- Week 7: Transparency and Explainability + *Guest talk* on Communicating DS Insights with Stakeholders by A/Prof Richard Xu, UTS
- Week 9: Human rights and wellbeing
- Week 10: AI/DS for Social Good (*Guest talk* by Prof Milind Tambe, Harvard University & Director, Google AI for Social Good)
- Week 11: Developing and Productising AI/DS project



Assignment Overview

—
What's next...

Recap: Assessment Structure



Task	Percentage	Due Date
Assessment Task 1 (Ethical AI/DS Case Study – Group) Case study Report Case study presentation	20%	Sundays 23:59 Week 7
Assessment Task 2 (AI/DS for Good Project – Group) Milestone 1 (M1 10%) – Initial Submission (Report) Milestone 2 (M2 40%) – Project Presentation and Report	50%	Sundays 23:59 M1: Week 5 M3: Week 11
Assessment Task 3 (Interviews & Reflections - Individual) Case study (Task 1) Reflection (5%) Project (Task 2) Reflection (5%) AI/DSP Interview (15%)	25%	Tuesdays 23:59 T1 reflection: Week 8 T2 reflection: Week 12 Interview: Week 14
Assessment Task 4 (Micro-credentials- Individual) Academic Integrity Awareness (1%) Presenting Using Story (3%) Ethical Cities (1%)	5%	Week 2 Week 2 (with badge submission due week 4) Week 9

FAQ



- **Micro-cred**
 - Just submit badge link – if done prior, resubmit the link to the badge
 - Content/assessment related Q: use the “ask your instructor a question” feature on Canvas
 - Technical problem: ask ITS
- **Group formation**
 - How many in a group? 3-4, not more than 4
 - Can group be formed across workshop? No
- **How is the project (Task 2) different from the Case Studies course?**
 - It's different because AI/DSP is research focused, no coding required
 - A strong focus on ethics and responsible AI/DS
 - Assessment around: Critical thinking, innovation, ethics, research components and methods, writing and presentation
 - An output of Task 2 is a solid research report and a project plan
- **Challenge Topics (broad) vs Case Study/Problem (specific)**
- **Research materials (references, data, evidence) must be easily accessible in the public domain**

Example from Past Assignment



Case study (Task 1)

- HD example
- D example

No example for Task 2

Assessment Timeline



- Week 2: Research Integrity; Presenting using Story micro-cred due
- Week 2-3: Group formation, topic discussion. By end of week 3: group to be registered on Canvas, and a channel to be set up on Teams
- Week 4: *Workshop* – discuss and finalise topic (Task 1 & Task 2) with tutor
- Week 5: Task 2 initial report due (after Easter break)
- Week 7: Task 1 (case study) due: report & recorded presentation
- Week 8: *Workshop* – Task 1 (case study) presentations (playback) & peer review; Task 1 individual reflection due
- Week 9: Ethical cities micro-cred due
- Week 11: Task 2 final report and presentation due
- Week 12: *Workshop* – Task 2 (project) presentations (playback) & peer review, Task 2 individual reflection due
- Week 14: AI/DSP individual interview

What to do this week?



- Micro-cred deadline reminder:
 - ✂ Micro-cred: Academic Integrity, Week 2
 - ✂ Micro-cred: Presenting using story, Week 2, link submission by Week 4
- Read materials
- Getting to know each other and form groups, next week is the deadline for notifying groups
- Think & discuss the topic for the project & case study (week 2 & 3)

Next Week



Guest lecture by Amazon AWS on Culture of Innovation – Working Backwards

Working with Data & Users

A brief run through of Research Methodology & Problem Formulation

Grand challenge presentation by Cybersecurity CIC & Communitier

Week 3 Guest talk: Culture of Innovation

(The talk won't be recorded)



"Amazon's approach to innovation has remained consistent since the company first launched - start with the customer and work backwards. Culture of Innovation provides valuable insights, lessons learned, and best practices from Amazon's cultural mechanisms by helping identify the next steps with the right AWS teams to solve problems and realize the opportunity for the end customers."

Catherine Eibner, Head of Cloud Innovation Programs from AWS will be giving the talk on Monday 15th.

Catherine's Bio: Catherine is a serial entrepreneur and a leader of innovation and growth in Australia's start-up economy. She built scalable programs and communities to support Entrepreneurs and Women in Technology while at Microsoft, then going on to become the GM of Start-ups at Incubator Blue Chilli. She joined AWS in late 2017, initially working in a start-up team in supporting and developing early stages of start-up companies and then recently joining the ANZ Public Sector business where she is the Digital Innovation lead, focused on improving the digital experiences for students and citizens by enabling digital transformation in the government, non-profit and education sectors through the innovative application of cloud technologies.

Catherine has been named one of the Top Australian Women in Tech and is recognised as one of the country's leading Start-up Mentors, actively advising start-ups, while also supporting entrepreneurs through her volunteer work and angel investments across the ecosystem.



Research integrity and ethics

A guest talk by Research Integrity and Governance team

Guest Talk: 'Research integrity and ethics'



This interactive session will introduce key concepts of responsible and ethical research. It includes discussion and activities focused on the principles and dilemmas of research integrity and human research ethics.

Dr Daniel Barr BSc(Hons) PhD *Melb*. Dan is the Principal Research Integrity Advisor at RMIT University. Dan is an experienced research manager specialising in how institutions enable research integrity. Dan also conducts research into research integrity. Dan is a Training Fellow of the Australasian Research Management Society and an Honorary Senior Fellow at the University of Melbourne.

Mary Duffy BEd(Sec)Arts, PGDipArts, MA *Melb*. Mary is one of the Coordinators, Research Governance and Ethics in the Research Ethics Integrity and Governance team, and is the coordinator for the STEM College Human Advisory Network (CHEAN). Mary has worked in research governance and ethics since 2009 and joined RMIT in 2016 from the University of Melbourne.

Dr David Blades BA(Hons) *UQ* PhD *RMIT*. David is Senior Coordinator, Research Integrity and Governance and Dr Daniel Barr's minion. He has a background in human research ethics and humanities, arts, and social sciences.

Dr Fotini Toso BA, PGDipArts, MA (Research), PhD. Fotini is one of the Coordinators, Research Governance and Ethics and is the coordinator for the DSC College Human Advisory Network (CHEAN). Fotini has worked in research governance and ethics since 2015 and joined RMIT in 2020 from the University of Divinity, where she managed the Human Research Ethics Committee.